

Philipp Gabler, BSc

Automatic Graph Tracking in Dynamic Probabilistic Programs via Source Transformations

Master's Thesis

to achieve the university degree of
Master of Science

submitted to
Graz University of Technology

Supervisor
Univ.-Prof. Dipl.-Ing. Dr. mont. Franz Pernkopf

Co-supervisor
Dipl.-Ing. Martin Trapp, BSc

Institute of Signal Processing and Speech Communication

Faculty of Electrical and Information Engineering

Graz, XXXX 2020

AFFIDAVIT

I declare that I have authored this thesis independently, that I have not used other than the declared sources/resources, and that I have explicitly indicated all material which has been quoted either literally or by content from the sources used. The text document uploaded to TUGRAZonline is identical to the present master's thesis.

Date

Signature

This work is licensed under a
[Creative Commons Attribution-ShareAlike 4.0 International License](https://creativecommons.org/licenses/by-sa/4.0/).



All code samples, unless otherwise noted or cited from other sources,
are also available under an [MIT license](#):

The MIT License (MIT)

Copyright (c) 2020 Philipp Gabler

Permission is hereby granted, free of charge, to any person obtaining a copy of this software and associated documentation files (the "Software"), to deal in the Software without restriction, including without limitation the rights to use, copy, modify, merge, publish, distribute, sublicense, and/or sell copies of the Software, and to permit persons to whom the Software is furnished to do so, subject to the following conditions:

The above copyright notice and this permission notice shall be included in all copies or substantial portions of the Software.

THE SOFTWARE IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO THE WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT. IN NO EVENT SHALL THE AUTHORS OR COPYRIGHT HOLDERS BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY, WHETHER IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING FROM, OUT OF OR IN CONNECTION WITH THE SOFTWARE OR THE USE OR OTHER DEALINGS IN THE SOFTWARE.

The \LaTeX source of this document is available at
<https://github.com/philpsgabler/master-thesis>
or upon request from the author.¹

¹pgabler@student.tugraz.at

ABSTRACT

This thesis presents a novel approach for the implementation of a tracking system to facilitate program analysis, based on program transformations. The system is then applied to a specific problem in the field of probabilistic programming.

The main contribution is a general system for the extraction of rich computation graphs in the Julia programming language, based on a transformation of the intermediate representation (IR) used by the compiler. These graphs contain the whole recursive structure of any Julia program in terms of executed IR instructions. The system is flexible enough to be used for multiple purposes that require dynamic program analysis or abstract interpretation, such as automatic differentiation or dependency analysis.

The second part of the thesis describes the application of this graph tracking system to probabilistic programs written for Turing, a probabilistic programming system implemented as an embedded language within Julia. Through this, an executed Turing model can be analyzed, and the dependency structure of involved random variables be extracted from it. Given this structure, analytical Gibbs conditionals can be calculated and passed to Turing's inference mechanism, where they are used in Markov-Chain Monte Carlo samplers approximating the modelled distribution.

Contents

Notation	xi
1 Introduction	1
1.1 Problem Description	1
1.2 Related Work	1
2 Background	3
2.1 Bayesian Inference and Probabilistic Programming	3
2.1.1 Probabilistic Programming	5
2.2 Computation Graphs and Automatic Differentiation	5
2.3 Metaprogramming and Compilation in Julia	5
3 Implementation of Dynamic Graph Tracking in Julia	7
3.1 Automatic Graph Tracking and Extended Wengert Lists	7
4 Graph Tracking in Probabilistic Models	9
4.1 Dependency Analysis in Dynamic Models	9
4.2 JAGS-Style Automatic Calculation of Gibbs Conditionals	9
4.3 Evaluation	9
5 Discussion	11
5.1 Future Work	11
Bibliography	13

Notation

$\mathbb{P}[\Theta \in A \mid X = x]$	Random variables and their realizations will usually be denoted with upper and lower case letters, respectively (with some exceptions for Greek variable names). Sets are written with uppercase letters.
$\mathbb{E}[X], \mathbb{V}_X[Y]$	Expectation and variance; if necessary, the variable with respect to which the moment is taken is indicated.
$f(x), \phi_Z(x)$	Density function are named using letters commonly used for functions, with an optional subscript indicating the random variable they belong to.
$p(x, y \mid z)$	The usual abuse of notation with the letter “p” standing for any density indicated by the names of the variables given to it is used as well (in this case, $f_{X,Y Z}$ is implied).
$\int p(x) dx = 1$	Integrals over the whole domain of a density are written as indefinite integrals, where the usage is clear.

1 Introduction

1.1 PROBLEM DESCRIPTION

1.2 RELATED WORK

2 Background

This section provides the background for the concepts used later in chapters 3 and 4. Initially, it gives a quick overview of Bayesian inference and probabilistic programming in general, necessary to understand the requirements and usual approaches of probabilistic programming systems.

Consequently, the machinery and language used to develop the graph tracking system forming the main part of the work are described. This consists firstly of a short introduction to graph tracking and source-to-source automatic differentiation, which contain many ideas and terminology that will be used later, and often provided inspiration. Secondly, the basic notions and techniques of the Julia compilation process as well as the language’s metaprogramming capabilities are described, which form the basis of the implementation.

2.1 BAYESIAN INFERENCE AND PROBABILISTIC PROGRAMMING

Probabilistic modelling is a technique for modelling phenomena based on the assumption that observables can be fully described through some stochastic process. When we assume this process to belong to a specified family of processes, the estimation of the “best” process is a form of learning: if we have a good description of how observations are generated, we can make summary statements about the whole population (descriptive statistics) or predictions about new observations. When observations come in pairs of independent and dependent variables, learning the conditional model of one given the other solves a regression or classification problem.

Within a Bayesian statistical framework, we assume that the family of processes used is specified by a joint distribution of random variables related through conditional distributions with densities, which describe how the observables would be generated: some *latent variables* are generated from *prior distributions*, and the observed *data* are generated conditionally on the latent variables. The goal is to learn the *posterior distribution* of the parameters given the observations, which is the inverse of how the problem is specified.

As an example, consider image classification: if we assume that certain percentages of an image data set picture cats and dogs, respectively, the distribution of these labels forms the prior. Given the information which animal is depicted on it, an image can then be generated as a matrix of pixels based on a distribution of images

conditioned on labels. The posterior distribution is then the “inverse” distribution of the label given an image. When we have this information, we can, for example, build a Bayesian classifier, by returning for a newly observed image that label which has the higher probability under the posterior.

This kind of learning is called Bayesian inference since, in the form of densities, the form of the model can be expressed using Bayes’ theorem:

$$p(\theta | x) = \frac{p(x | \theta) p(\theta)}{p(x)}, \quad (2.1)$$

where x is the vector of observed data, and θ are the parameters.¹ (In practice, not all of the latent variables have to parameters we are actually interested in; these can be integrated out).

Going beyond simple applications like the classifier mentioned above, using the posterior gets difficult, though. Simply evaluating the posterior density $\theta \mapsto p(\theta|x)$ at single points is not enough for usages such as parameter estimation with continuous parameters, or even simple probability evaluation; and even for discrete variables, combinatorial explosion limits the usability of the simple summation approach to calculate probabilities. The reason is that almost all of the relevant quantities depend on some sort of expectation over the posterior density, which is an integral:

$$\mathbb{E}[f(\Theta) | X = x] = \int f(\theta)p(\theta | x) d\theta, \quad (2.2)$$

with the base measure depending on the type of Θ . When the distributions involved form a sufficiently “nice” combination, e.g., when they form a conjugate pair, the integrals can be calculated analytically since the posterior density has a closed form for a certain known distribution, or at least is a known integral. In general, however, this is not possible, and approximation have to be made, for which different approaches are available: for example, expectation maximization for models with latent variables or message passing and variational methods for general graphical models.

The applications described in this thesis fall into another framework for approximating posterior information: Monte Carlo methods, which derive from a specified density a procedure that provides a sample of random variables for which a law of large numbers (LLN) holds:

$$\frac{1}{n} \sum_{i=1}^n f(\Theta_i) \rightarrow \mathbb{E}[f(\Theta) | X = x], \quad \text{as } n \rightarrow \infty \quad (2.3)$$

in some appropriate stochastic convergence (usually convergence in probability is enough). When $\Theta_i \sim p(\cdot | x)$ can be sampled exactly, and the Θ_i are thus i.i.d., the LLN holds trivially; such samplers exist, but might also be difficult to derive or not possess good enough convergence properties. Another large class of samplers is

¹Note the abuse of notation with $p()$ and the integral; see page xi on notation.

formed by *Markov Chain Monte Carlo* (MCMC) methods, which, instead of sampling exactly from the density, provide a Markov chain whose stationary distribution is the target density. The advantage of MCMC methods is that they mostly treat the density as a black box, without requiring any more formal manipulations on it.

The general scheme of MCMC methods consists of the definition of a transition kernel by means of two helper functions: a proposal distribution with density q , and an acceptance rate α , both depending on the old value in the chain:

1. Start from an arbitrary $\Theta_1 = \theta_1$.
2. For each $i \geq 1$:
 1. Propose $\hat{\Theta}_i \sim q(\cdot \mid \Theta_{i-1})$.
 2. With probability $\alpha(\hat{\Theta}_i, \Theta_{i-1})$, set $\Theta_i = \hat{\Theta}_i$; else, keep $\Theta_i = \Theta_{i-1}$.

2.1.1 Probabilistic Programming

Probabilistic programming is a means of describing probabilistic models through the syntax of a programming language. Probabilistic programs distinguish themselves from normal programs by the possibility of being sampled from conditionally, with some of the internal variables fixed to observed values. While probabilistic programming systems are often implemented as separate, domain-specific languages, they can also be embedded into “host” programming languages with sufficient syntactic flexibility. The latter is advantageous if one wants to use regular general-purpose programming constructs or interact with other functionalities of the host language.

2.2 COMPUTATION GRAPHS AND AUTOMATIC DIFFERENTIATION

2.3 METAPROGRAMMING AND COMPILATION IN JULIA

3 Implementation of Dynamic Graph Tracking in Julia

3.1 AUTOMATIC GRAPH TRACKING AND EXTENDED WENGERT LISTS

4 Graph Tracking in Probabilistic Models

4.1 DEPENDENCY ANALYSIS IN DYNAMIC MODELS

4.2 JAGS-STYLE AUTOMATIC CALCULATION OF GIBBS CONDITIONALS

4.3 EVALUATION

5 Discussion

5.1 FUTURE WORK

COLOPHON

This document was typeset using the pdf_{La}TeX typesetting system, with the memoir document class. The body text is set in 11 pt Linux Libertine, enhanced by the microtype package. Other fonts include Biolinum and Inconsolata.

The document source has been written in Emacs with AU_CTeX mode, using TeXworks as PDF viewer.